

Artificial Intelligence and Machine Learning Applications

Lecture Outline

- Linear Regression Tutorial & Playground
- Image Compression
- Logistic Regression Interactive Demo
- Neural Network Regression Tutorial & Playground
- Auto-Encoders

Linear Regression Tutorial



<https://mlu-explain.github.io/linear-regression/>

Linear Regression Playground

<https://observablehq.com/@yizhe-ang/interactive-visualization-of-linear-regression>



Image Compression with Linear Regression

- An image can be transformed into the frequency domain and represented as a combination of some basic components.
- Cosine Basis Functions and Discrete Cosine Transforms (DCT) can enable this.
- The human eye is most sensitive to low frequencies. Therefore, most of the high frequencies can be ignored.
- Principal behind JPEG Image Compression.

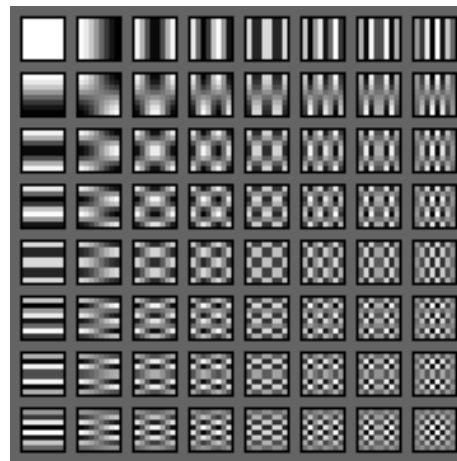



Image Compression with Linear Regression

- Using linear regression, we can learn the weight of each of these cosine basis.
- We only need to store the weights of the basis frequencies.
- The image is reconstructed using these weights.

Image Compression

8 x 8 Pixels



Image



Image Compression

- Gray-Scale Example:
- Value Range 0 (black) --- 255 (white)

| | | | | | | | |
|-----|-----|-----|-----|-----|----|-----|-----|
| 63 | 33 | 36 | 28 | 63 | 81 | 86 | 98 |
| 27 | 18 | 17 | 11 | 22 | 48 | 104 | 108 |
| 72 | 52 | 28 | 15 | 17 | 16 | 47 | 77 |
| 132 | 100 | 56 | 19 | 10 | 9 | 21 | 55 |
| 187 | 186 | 166 | 88 | 13 | 34 | 43 | 51 |
| 184 | 203 | 199 | 177 | 82 | 44 | 97 | 73 |
| 211 | 214 | 208 | 198 | 134 | 52 | 78 | 83 |
| 211 | 210 | 203 | 191 | 133 | 79 | 74 | 86 |




Image Compression

- 2D-DCT of matrix

Numbers are coefficients
of polynomial

```
-304 210 104 -69 10 20 -12 7
-327 -260 67 70 -10 -15 21 8
 93 -84 -66 16 24 -2 -5 9
 89 33 -19 -20 -26 21 -3 0
 -9 42 18 27 -7 -17 29 -7
 -5 15 -10 17 32 -15 -4 7
 10 3 -12 -1 2 3 -2 -3
 12 30 0 -3 -3 -6 12 -1
```


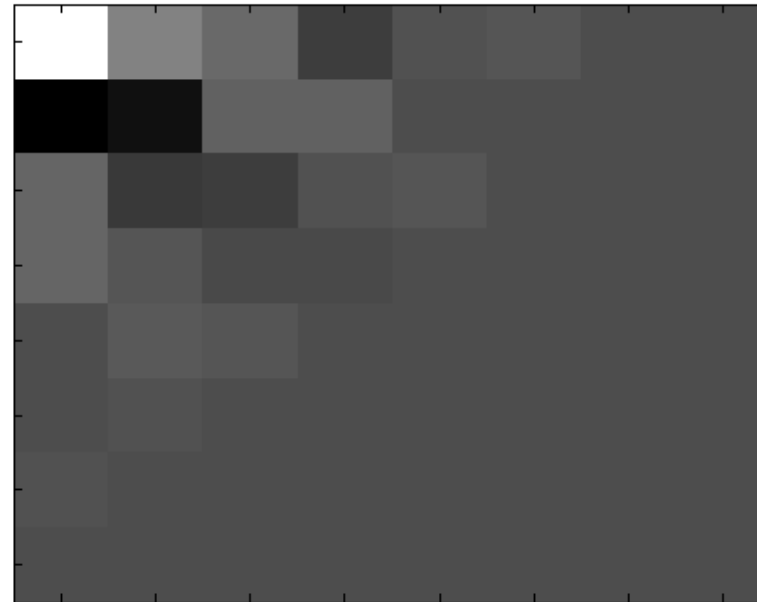


Image Compression

- Cut the least significant components

```
-304 210 104 -69 10 20 -12 0
-327 -260 67 70 -10 -15 0 0
 93 -84 -66 16 24 0 0 0
 89 33 -19 -20 0 0 0 0
-9 42 18 0 0 0 0 0
-5 15 0 0 0 0 0 0
10 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0
```




As you can see, we save a little over half the original memory.

Reconstructing the Image


- New Matrix and Compressed Image

| | | | | | | | |
|-----|-----|-----|-----|-----|----|----|-----|
| 55 | 41 | 27 | 39 | 56 | 69 | 92 | 106 |
| 35 | 22 | 7 | 16 | 35 | 59 | 88 | 101 |
| 65 | 49 | 21 | 5 | 6 | 28 | 62 | 73 |
| 130 | 114 | 75 | 28 | -7 | -1 | 33 | 46 |
| 180 | 175 | 148 | 95 | 33 | 16 | 45 | 59 |
| 200 | 206 | 203 | 165 | 92 | 55 | 71 | 82 |
| 205 | 207 | 214 | 193 | 121 | 70 | 75 | 83 |
| 214 | 205 | 209 | 196 | 129 | 75 | 78 | 85 |



Can You Tell the Difference?

Original



Compressed




Image Compression

Original



Compressed



Logistic Regression Interactive Tutorial



<https://mlu-explain.github.io/logistic-regression/>

Neural Networks Tutorial



<https://mlu-explain.github.io/neural-networks/>

Neural Networks Playground



<https://playground.tensorflow.org/>

Try it Yourself – Digit Classification

<https://trekhleb.dev/machine-learning-experiments/#/experiments/DigitsRecognitionMLP>


Try it Yourself – Sketch Recognition

<https://trekhleb.dev/machine-learning-experiments/#/experiments/SketchRecognitionMLP>

Applications – AutoEncoders

- Autoencoders are a type of neural networks where the input is also the output.
- They come under unsupervised learning and there are no labels involved.
- An autoencoder consists of two parts: encoder and decoder.
- The idea here is that you take a higher dimensional input, project it into a lower dimensional space and then project it back into the input space.

Applications – AutoEncoders



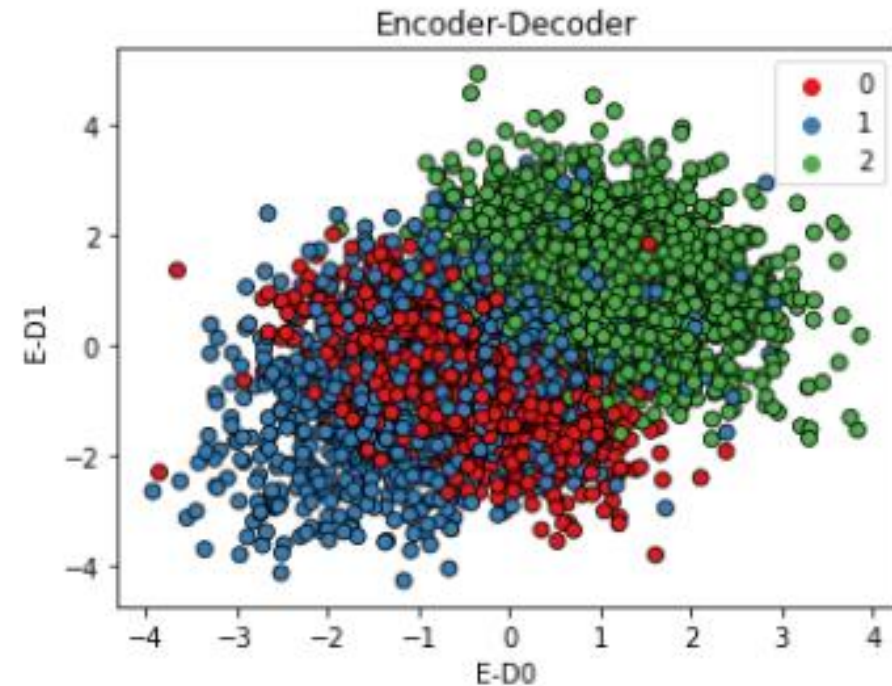
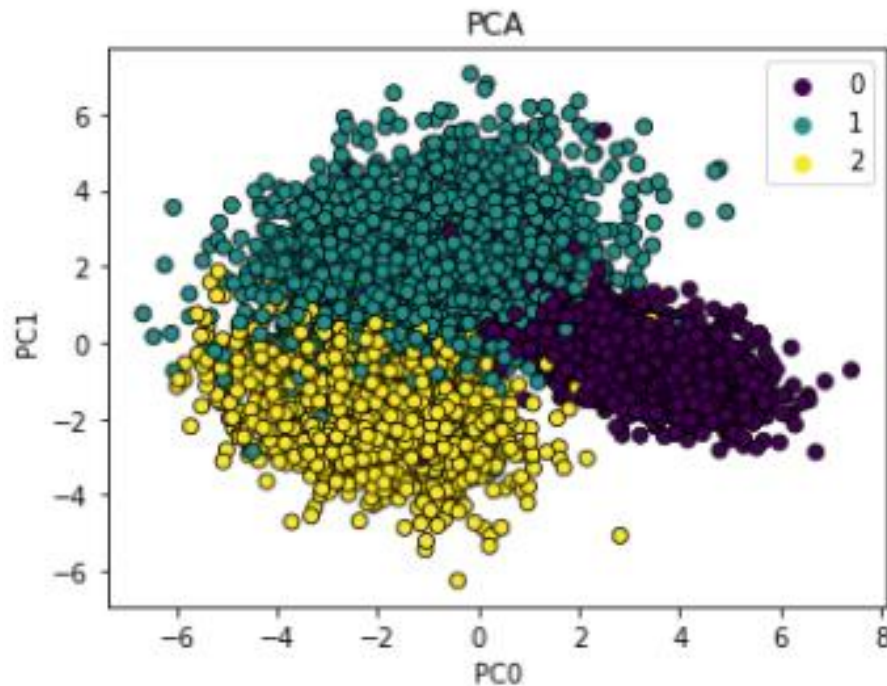
Applications – AutoEncoders

- The autoencoder model tries to minimize the reconstruction error (RE).
- Typically, mean squared is used as the loss function for autoencoders.
- The objective is to minimize the following:

$$L(x, \hat{x}) = \frac{1}{N} \sum_{i=1}^N ||x_i - \hat{x}_i||^2$$


AutoEncoders for Dimensionality Reduction

- The encoder output can be used for dimensionality reduction

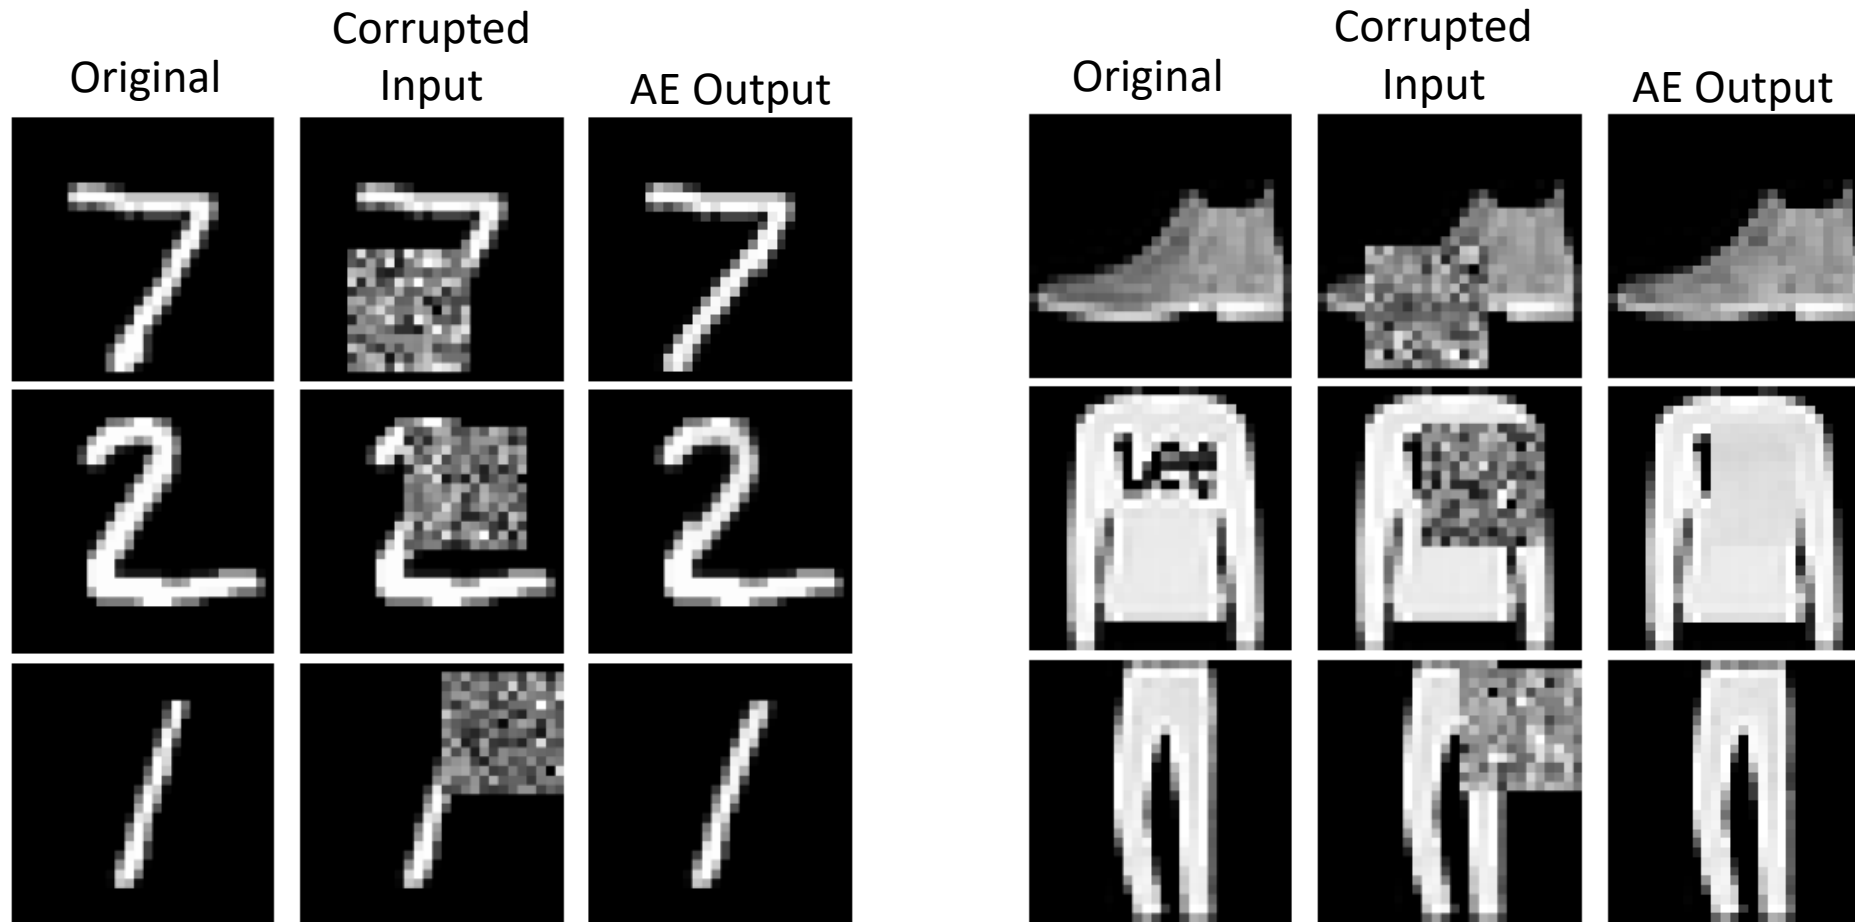


AutoEncoders for Outlier Detection

- The reconstruction error can be used to detect outliers. Out of distribution samples will have high Reconstruction error.



AutoEncoders for Image Completion



Summary

